

What the 2026 AFL draw was worth

A method note on separating list quality from fixture quality.

Gaurav Warke · September 2026 · Season 2026, 2,291 matches

This note exists so the numbers can be checked by someone who was not in the room when they were presented. It states what was estimated, what came back significant, what did not, and where the method stops being reliable. The nulls are reported alongside the findings, because a method that only reports what it found is not a method.

The question

Two clubs finish a season four places apart. How much of that gap is the players and how much is the fixture? Every club argues this internally in November. The argument is usually settled by whoever is most senior in the room, because nobody has a number.

This estimates the number.

Data and model

Every men's AFL match from 2015 to 2026, 2,291 games, from the AFL's own match records via the `fitzRoy` R package. The outcome modelled is **match margin in points**. The fixture variables are attached to each team-match and are known before the game is played, so none of them can absorb the result they are predicting.

A linear model on margin, with team and opponent strength controlled, leaves four fixture effects to estimate:

Effect	Points of margin	95% CI	p Verdict
Playing on your own ground	+1.75	-0.56 to +4.06	0.138 not significant
Having travelled interstate	-5.80	-9.08 to -2.53	0.0005 significant
Each hour of time-zone change	-1.50	-3.52 to +0.53	0.147 not significant
Coming off a six-day break	+0.90	-1.42 to +3.21	0.449 not significant

Model $R^2 = 0.446$ on $n = 2,291$.

One of these four is real. Only the travel term is applied in everything downstream. The other three are reported and then set to zero, rather than kept in at a small value because the sign looked plausible.

The three findings, in order of how well they hold up

1. The six-day break is worth nothing

+0.9 points, confidence interval spanning zero in both directions, $p = 0.45$. On this data a club coming off five days' rest and a club coming off nine are indistinguishable once you know who they are playing and where.

This is the most robust claim in the note because it is a null. It does not depend on a modelling choice going a particular way; the effect simply is not large enough to detect in 2,291 games. If a six-day break were worth even two goals, this sample would see it.

2. Home ground advantage is mostly not the ground

Playing at your own venue is worth +1.75 points and is not statistically distinguishable from zero. Having flown interstate to get to a match costs 5.8 points, and that one is solid ($p = 0.0005$).

Home ground advantage in the AFL is therefore better described as an *away* disadvantage. It is not familiarity with the ground, the crowd, or the goal posts. It is that the other side got on a plane. A Melbourne club hosting another Melbourne club at the MCG has almost no edge at all, and that is exactly what the data shows.

3. The two Perth clubs are the biggest winners from travel

A trip to Perth costs a visiting side **8.8 points** — the most expensive journey in the competition.

West Coast and Fremantle each made 10 interstate trips in 2026, more than any Victorian club. They still finish first and third on the travel ledger:

Club	Net travel effect	Interstate trips
West Coast	+20.7	10
Port Adelaide	+18.8	10
Fremantle	+17.4	10
Adelaide	+11.7	10
...		
North Melbourne	−23.7	8
Hawthorn	−28.5	10

The reason is one-sided accounting. The Perth clubs pay the travel cost on their own trips, but every visitor to Perth pays the largest penalty on the schedule. They fly a great deal and still come out ahead, because the flight they impose on others is worse than the flight they take themselves.

A caveat that belongs with this finding: the *net* season effect for both Perth clubs has a confidence interval crossing zero (West Coast +23.6, CI −3.0 to +48.6; Fremantle +9.7, CI −23.8 to +49.3). The direction is consistent and the mechanism is clear, but the total is not individually significant for either club. The travel term underneath it is.

The three ladders

Three ladders are produced, and the distinction matters more than any single position.

Ladder	What is removed	Answers
Actual	nothing	What happened
Earned	goal-kicking luck	What their football won
Level field	kicking luck <i>and</i> the fixture	What the list is worth

“Earned” replaces each match result with the scoring-shot differential, on the basis that creating shots repeats from one half of a season to the next (reliability 0.88) while converting them barely does (0.25). “Level field” then removes the fixture burden on top.

The level-field ladder is the one to quote. It is the complete adjustment. The earned ladder is an intermediate step

shown for transparency, and a club will sometimes sit a place apart on the two.

Across the competition the draw was worth a spread of **135.8 points of margin** between the most and least favoured club — an average of 5.9 points a game.

Limitations

The level-field ladder is a counterfactual, not a result. It says what the ladder would look like if every club had faced an identical schedule. No club played that season. Treat it as a way of comparing lists, not as a claim that a different team deserved the flag.

The travel coefficient is an average. It does not distinguish a Thursday night flight from a Saturday morning one, a charter from a commercial service, or a club with a settled interstate routine from one without. Clubs hold the data that would separate these; this note does not.

Opponent strength is measured within the same season. Strength of schedule is therefore partly circular: a club looks to have had a hard draw partly because the sides it played won games, some of which were against that club.

The prediction model is good, not remarkable. Walking forward through the 2026 season it scores a log-loss of 0.5229 over 217 matches, against 0.693 for a coin toss, and tips at 69.4%. It beat six alternatives, but only three of those gaps are statistically real: ridge regression ($p = 0.18$) and Elo ($p = 0.28$) are tied with it on this sample.

And it has been overconfident about premierships. Replaying 2016 to 2025 under the old finals system, the side it installed as favourite won 2 flags out of 10, against 3.7 expected. That is not damning on ten observations ($P(\leq 2) = 0.22$) but it points one way, and September is the part of the year this model finds hardest.

Nothing here is causal. These are associations measured on observational data. The fixture is not randomly assigned: the AFL schedules blockbusters, protects broadcast windows, and gives travel concessions. Some of what is attributed to travel may be scheduling that correlates with it.

Reproducing this

All figures come from `outputs/fixture_equity.json` and `outputs/level_field.json`, produced by `app/build/fixture_audit.py` and `app/build/level_field.py`. Data is refreshed from the AFL via `R/06_live_data.R`. The model bake-off that selected the prediction engine is documented separately in `docs/model-bakeoff.md`.

Figures in this note were generated on 21 September 2026 from data current to the 2026 preliminary finals.

Questions, corrections and arguments are all welcome. The corrections are the most useful.